

Inverse Complexity Scaling: How Distributed Orchestration and Constitutional Autonomous Infrastructure Become More Stable, Efficient, and Cost-Effective as Complexity Increases

Leon Calvin Long II

DSOS-CAI Research — Independent Systems Laboratory, Spicewood, TX, USA

primary contributions: (1) a formal mathematical definition of ICS grounded in performance vector analysis, (2) a first-principles derivation showing ICS emerges from redundancy multiplication, route order-statistics, agent load distribution, and fixed-cost amortization, (3) an architectural explanation of the DSOS and CAI mechanisms that enable ICS, and (4) receipted operational evidence from a production DSOS/CAI substrate — zero false positives across a 200-row incident annex, deterministic confidence gating, and a 267× orchestration benchmark speedup — together with model-derived scaling projections for multi-scale deployments. The formal model projects up to 18× cost reduction, 28× throughput improvement, and 15× variance reduction between modeled small- and large-scale configurations; these projections are explicitly labeled as unmeasured and are proposed for independent validation via a community benchmark suite (ICS-Bench). We argue ICS fundamentally inverts classical infrastructure sizing heuristics for self-organizing, agent-parallel workloads and propose a community benchmark suite (ICS-Bench) for independent validation.

1. Introduction The dominant intellectual framework governing distributed systems design holds that complexity is an adversary. This view is codified across foundational texts: Brooks's No Silver Bullet [1] identifies complexity as the central difficulty of software engineering; Tanenbaum [2] catalogs the communication overhead that grows with node count; Amdahl's Law [3] places a hard ceiling on parallelization gains. The practical engineering maxim derived from these insights is explicit: start small, keep it simple, scale only under duress. This paper reports a result that contradicts this heuristic in a specific but increasingly important class of systems. In the author's production DSOS and CAI substrate and its chaos-injection validation corpus, a consistent pattern emerged: a self-organizing, role-differentiated, deterministically governed architecture absorbed failures with zero false-positive heals across the full receipted incident annex, while the formal model developed in this paper predicts that stability, determinism, throughput, and cost-efficiency all improve as node, route, and agent counts grow. The addition of complexity, rather than introducing overhead, introduces redundancy, load distribution, and amortization. Receipted operational evidence and model-derived projections are kept rigorously separate throughout this paper. This phenomenon—which we term Inverse Complexity Scaling (ICS) —is not a minor performance curiosity. It has substantive implications for how self-organizing, agent-parallel systems should be designed, sized, and costed. If ICS is real and

reproducible, then the standard engineering intuitions about right-sizing distributed AI infrastructure are systematically wrong for this class of architecture. The mechanism behind ICS is, in retrospect, comprehensible from first principles. In classical systems with fixed topologies and centralized coordinators, complexity yields primarily overhead: more components mean more failure modes, more communication, more coordination cost. In self-organizing systems with autonomous rerouting, role-differentiated agents, and shared

context graphs, additional complexity yields primarily redundancy, load distribution, and amortization—forces that drive performance upward rather than downward. The distinction between these two regimes has not been formally characterized in the literature. This paper provides that characterization.

1.1 Contributions This paper makes the following contributions:

1. Formal Definition. A rigorous mathematical definition of ICS as a property of performance vectors over a complexity dimension κ , with explicit monotonicity conditions for stability, throughput, determinism, and cost-efficiency.
2. Mathematical Derivation. First-principles derivations showing how route redundancy, agent load distribution, order-statistics of path selection, and fixed-cost amortization each contribute to ICS behavior.
3. Architectural Explanation. A detailed account of the DSOS and CAI architectural mechanisms—self-healing route tables, backpressure propagation, role-differentiated agents, shared context DAGs, and lightweight local consensus—that together enable ICS.
4. Receipted Evidence and Model Projections. Operational results from the production substrate's incident ledger (stability, governance, and benchmark metrics), plus model-derived multi-scale projections explicitly labeled as unmeasured.
5. Design Implications. A revised infrastructure design philosophy and cost model appropriate for ICS-exhibiting systems.

6. Community Roadmap. A proposal for ICS-Bench, a formal benchmark suite for independent community validation.

2. Background and Related Work

2.1 Classical Scaling Models The theoretical underpinnings of distributed systems performance have been shaped by several landmark results. Amdahl's Law [3] establishes that the speedup from parallelization is bounded by the serial fraction of a workload, implying diminishing returns as node count grows. While strictly about parallelization, Amdahl's formulation has been widely generalized to argue against complexity scaling. Gustafson's Law [4] offers a more optimistic counter-argument for scaled problem sizes but does not address fault behavior or cost. Conway's Law [5] argues that system architectures mirror the communication structures of the organizations that produce them—implying that complexity in organizational coupling translates to complexity in system coupling. This sociological observation reinforces the intuition that complex systems are difficult to reason about and expensive to maintain. Communication overhead theory (captured in Metcalfe's Law and its variants) holds that network communication cost grows as $O(N^2)$ with the number of participants N . For systems requiring all-pairs coordination, this is a hard scaling wall. DSOS and CAI, as we demonstrate, avoid all-pairs coordination by design.

2.2 Fault Tolerance and Redundancy Theory The literature on fault-tolerant systems has long recognized that redundancy improves reliability. Reliability theory establishes that for independent components with individual failure probability p , the probability that all N components fail simultaneously is p^N , an exponentially declining function [6]. This result is foundational to RAID storage, quorum-based consensus, and multi-path networking. What has not been previously explored is how this reliability benefit interacts with throughput, latency determinism, and cost-efficiency simultaneously,

and whether a system can be designed such that all four metrics improve together as a function of a unified complexity parameter. Lamport's work on distributed consensus [7] and the Raft consensus algorithm [8] (Ongaro & Ousterhout) provide the theoretical and practical basis for the lightweight local consensus mechanisms employed in DSOS. Raft's key insight—that leader election and log replication can be made comprehensible without sacrificing correctness—informs DSOS's approach to route table consensus.

2.3 Complex Adaptive Systems and Emergence Axelrod and Cohen [9] demonstrated that complex

adaptive systems can exhibit emergent properties that are not predictable from analysis of individual components. Ant-colony optimization [10] and gossip protocols [11] provide concrete computational instances of emergence: globally useful behaviors (shortest-path routing, epidemic information spread) arising from locally simple agent rules without central coordination. DSOS's autonomous rerouting and CAI's agent collaboration draw directly on these principles.

2.4 Multi-Agent Systems and CAI-Adjacent Literature

Multi-agent systems (MAS) research [12] has extensively studied coordination, negotiation, and task allocation among autonomous agents. However, the MAS literature has predominantly focused on correctness and completeness of coordination protocols rather than on how performance metrics scale with agent count. The CAI framework introduced in this paper extends MAS theory with an explicit account of role differentiation, backpressure signaling, and shared context graphs as mechanisms for ICS.

2.5 Why Existing Models Fail to Capture ICS

Existing scaling models universally assume one or more of the following: (a) fixed topology communication graphs, (b) centralized or hierarchical coordinators, or (c) homogeneous node/agent populations. Each assumption is violated by DSOS and CAI. Fixed topologies cannot exploit route diversity; centralized coordinators create bottlenecks that worsen with scale; homogeneous populations cannot benefit from

role-differentiated load distribution. The ICS phenomenon is precisely what emerges when all three assumptions are relaxed simultaneously.

3. System Architecture Overview

3.1 DSOS: Distributed Self-Organizing Systems

A Distributed Self-Organizing System (DSOS) is a graph-structured orchestration layer in which nodes are autonomous execution units that communicate via typed, versioned message channels. Unlike conventional service meshes, DSOS nodes maintain local route tables that are continuously updated through lightweight consensus, enabling the system to heal routes and rebalance load without any central coordinator. Formally, a DSOS is defined as a directed graph $G = (V, E)$ where each vertex $v \in V$ is a node instance and each edge $e \in E$ is a typed, versioned channel with associated health metadata. The graph topology is dynamic: edges are created, retired, and rerouted autonomously by node-level Sentinel processes.

Key properties of DSOS:

- **Self-healing route tables:** Each node maintains a local route table updated via Raft-variant consensus with immediate neighbors. Route failures are detected within a configurable heartbeat interval and failover paths are precached.
- **Local consensus, no global coordinator:** Consistency is achieved through overlapping local quorums rather than a global leader, eliminating single points of coordination failure.
- **Typed, versioned channels:** Message schemas carry version tags, enabling backward-compatible evolution of message formats without routing disruption.

DSOS Node Types:

- **Ingress Node:** Accepts external requests; performs validation, schema normalization, and initial routing decisions.

- **Processor Node:** Executes computational workloads; stateless by design to enable horizontal scaling.
 - **Router Node:** Maintains route tables; makes forwarding decisions based on channel health and load signals.
 - **Egress Node:** Aggregates results and delivers responses to external consumers.
 - **Sentinel Node:** Continuously probes route health, detects anomalies, precaches failover paths, and updates neighbor route tables.
- [Figure 1. DSOS node topology at illustrative medium scale (32 nodes): Ingress, Router, Processor, Egress, and Sentinel node types interconnected by typed, versioned channels with local-neighborhood consensus. Diagram omitted in this text rendering.]

the probability of complete routing failure is: $P(\text{all routes fail}) = (p^N)^R = p^{(N \cdot R)}$

Since both N and R grow with complexity κ , the failure probability decreases superexponentially in κ . Correspondingly, MTBF (the reciprocal of failure rate) grows at least as: $\sigma(\kappa) = \Omega(\log \kappa)$ The receipted incident record (Section 6.2) is consistent with this super-exponential absorption within the production substrate; whether MTBF grows strictly faster than logarithmically across scales requires the multi-scale measurement grid of Section 6.4 and ICS-Bench. 4.2.2 Throughput: Route Diversity and Load Distribution With R available routes to a destination, each with latency drawn independently from distribution $F(l)$, a Router node that selects the minimum-latency available route observes an expected latency equal to the expected minimum order statistic: $E[L_{\min}] = E[\min(L_1, L_2, \dots, L_R)]$ By order statistics theory, as $R \rightarrow \infty$, $E[L_{\min}]$ converges to the essential infimum of F (the minimum achievable latency). Thus route diversity drives expected latency toward the physical minimum, enabling throughput to grow as: $\tau(\kappa) = \Omega(\kappa^\alpha)$ for some $\alpha \in (0, 1)$ The sub-linear exponent $\alpha < 1$ reflects that throughput gains are real but bounded by the physical channel capacities; the inequality $\alpha > 0$ confirms that gains persist across the ICS operational range. For agent load distribution with M Executor agents processing a Poisson arrival stream of rate λ , per-agent queue depth is approximately: $Q_i \approx \lambda / M$ As M increases (more Executor agents), per-agent contention decreases, reducing queueing delay and increasing effective system throughput. 4.2.3 Determinism: Variance Reduction The variance of system output latency is dominated by the variance of the maximum latency component in the execution path. With multiple parallel Executor agents

handling independent task lanes, and the system taking the result of whichever lane completes first, the effective latency variance follows: $\text{Var}[L_{\text{effective}}] = O(1 / \kappa)$ This result follows from the central limit theorem applied to the distribution of minimum-order statistics: as the number of parallel paths increases, the variance of the minimum converges at rate $O(1/N)$. Formally: $\delta(\kappa) = O(1 / \kappa)$ 4.2.4 Cost-Efficiency: Fixed-Cost Amortization Total system cost has two components: fixed infrastructure overhead C_{fixed} (coordination layer, Sentinel probes, shared context graph maintenance) and variable cost C_{var} per operation. Cost-per-operation is thus: $\rho(\kappa) = C_{\text{fixed}} / \text{Ops}(\kappa) + C_{\text{var}}$ Since $\text{Ops}(\kappa) = \Omega(\kappa^\alpha)$ grows with κ , the fixed-cost term decreases as κ increases: $\rho(\kappa) = O(C_{\text{fixed}} / \kappa^\alpha + C_{\text{var}})$ At sufficient scale, C_{var} also decreases because shared context graph reuse reduces per-operation memory allocation: if context memory cost C_{context} is fixed and shared across M agents, per-agent memory cost is C_{context} / M . The combined effect produces the cost-amortization regime projected in Section 6.4. 4.3 Phase Transition Behavior ICS is not a globally linear phenomenon. It exhibits distinct behavioral phases as κ varies across the operational range. Below $\kappa_{\min}$, the system is fragile: too few routes exist for redundancy to provide meaningful protection, agent count is insufficient for meaningful load distribution, and the fixed overhead of the DSOS/CAI coordination layer dominates costs. Performance in this regime is worse than a simple monolithic system. Within the range $[\kappa_{\min}, \kappa_{\text{inflection}}]$, ICS benefits accumulate rapidly. Each additional node or agent contributes disproportionately to redundancy and load distribution because

the marginal value of an additional route or agent is highest when the baseline count is low. This is the superlinear improvement zone. Within $[\kappa_{\text{inflection}}, \kappa_{\max}]$, ICS benefits continue at a diminishing but still positive rate. The system has entered its high-performance steady state; additional complexity provides incremental gains. Beyond $\kappa_{\max}$, coordination overhead from Raft consensus rounds, Sentinel probe traffic, and context graph synchronization begins to dominate. The system transitions to a classical complexity-penalty regime. Determining $\kappa_{\max}$ empirically is a critical research priority (see Section 7.4).

shared context graph, execute the task, and write outcomes back. This statelessness enables horizontal scaling without state migration: adding an Executor agent requires only registering it with the Planner's agent pool. As Executor count grows, per-agent queue depth falls as $Q_i \approx \lambda / M$, reducing both latency and variance. Reflector agents introduce a form of online learning: they analyze execution outcomes, identify systematic inefficiencies (e.g., task decompositions that consistently produce rework), and update Planner behavioral policies. This means system efficiency improves with runtime independently of any offline training, contributing to a time-dimension component of ICS beyond purely structural scaling.

5.4 Shared Context Graphs and Memory Efficiency The CAI context graph is a versioned, append-only DAG maintained in shared memory accessible to all agents. Because the graph is shared rather than replicated per agent, the memory footprint of context is $O(1)$ with respect to agent count, rather than $O(M)$ as it would be in a copy-per-agent architecture. More precisely: if C_{context} is the fixed memory cost of the current context graph version, and this cost is shared across M agents, then the effective per-agent memory cost for context is C_{context} / M . As M increases, each agent "pays" a smaller share of the context overhead. This is a direct driver of the cost-amortization projections of Section 6.4. The append-only, versioned structure of the DAG also enables lock-free concurrent reads: agents read specific version snapshots without blocking each other. Write conflicts are impossible because all writes are append operations to new version nodes. This eliminates lock contention as a scaling bottleneck, allowing read throughput to scale linearly with agent count.

5.5 Self-Healing Consensus Mechanisms DSOS nodes maintain consistent route table state through a Raft-variant lightweight consensus protocol operating over local neighborhoods (each node participates in consensus only with its immediate neighbors and their neighbors, not with the entire system). This locality constraint bounds consensus communication cost at $O(1)$ per node, independent of total system size. Quorum-based consensus requires that more than $N/2$ nodes in a quorum group remain live for the group to make progress. The probability of quorum loss—requiring a majority of quorum-group nodes to fail simultaneously—decreases exponentially with quorum group size: $P(\text{quorum loss}) = P(\text{more than } N/2 \text{ nodes fail}) = \sum_{k=\lfloor N/2 \rfloor + 1..N} C(N,k) \cdot p^k \cdot (1-p)^{N-k} \rightarrow 0$ exponentially in N . As system complexity κ increases, quorum groups grow (or more groups are formed with overlapping membership), compounding this exponential reliability benefit. The local-neighborhood restriction ensures that the consensus mechanism itself does not

become an $O(N^2)$ communication bottleneck—a critical design decision that preserves ICS scaling.

6. Evidence from the Production DSOS/CAI Substrate

Amendment note (Version 2). Version 1 of this section presented model-derived figures (MTBF, throughput, latency variance, and cost at 8/32/128-node scales) as measured results from multi-scale production deployments. No such multi-scale measurement grid was operated; those figures are projections of the formal model. Version 2 reports only what the receipts support — operational results from the author's production substrate — and retains the multi-scale figures below as explicitly labeled projections. This amendment is disclosed rather than silently applied, consistent with the receipt-first doctrine under which the underlying substrate is operated.

6.1 Substrate Configuration and Evidence Provenance

The receipted evidence in this paper comes from a single production DSOS/CAI substrate: a 25-node, 10-agent deployment that orchestrates a family of distributed applications under bifurcated governance — probabilistic agents propose actions, and a deterministic runtime layer validates every action against frozen invariant contracts before execution. Stability and governance evidence is drawn from a 200-row incident annex (LVR-0907) spanning the substrate's chaos-injection validation corpus and production healing record, together with a 400-injection chaos run executed against the live substrate. Healing is playbook-driven and threshold-frozen: an anomaly is auto-healed only when its detection confidence meets the per-playbook frozen threshold (e.g., 0.70 for execute-class playbooks, 0.85 for critical-approval-class playbooks); otherwise it fails closed to human review. All records are server-timestamped and published in a versioned public incident ledger, providing an auditable provenance chain from detection to resolution.

6.2 Receipted Stability and Governance Results

Across the 200-row incident annex: 188 anomalies were resolved by playbook-driven auto-healing; 12 were confidence-gated (detected, scored below their frozen per-playbook thresholds, and escalated to human review — never auto-healed); zero false-positive heals occurred; and zero sub-threshold auto-executions occurred. Every escalated critical followed the human-in-the-loop path — no autonomous healing of critical flows.

On the 400-injection chaos run executed against the live substrate, the auto-heal ratio was 84.5% versus an 88.5% historical baseline, with zero false positives; all sub-threshold anomalies failed closed to human review rather than being retried or force-healed. A preregistered temporal-standing test (treatment: a mid-flight re-score of a detected anomaly dropping below threshold; control: an identical anomaly without re-score) confirmed the fail-closed refusal property: the treatment case refused healing after the fresh score fell below threshold, while the control executed normally. The falsifier specification was published before execution and the evidence rows after, in the public ledger.

These receipted results demonstrate the ICS-relevant properties measurable at a single scale: deterministic failure absorption, threshold-frozen governance with no heuristic drift at runtime, zero false-positive healing under 400-injection load, and monotone respect of confidence gates under mid-flight re-scoring.

6.3 Receipted Orchestration Benchmark

The Dual Mesh benchmark — the substrate's standard orchestration benchmark (25 active nodes, 10 agents) — measures dual-mesh operation (probabilistic proposal mesh plus deterministic validation mesh) against the single-mesh baseline. The receipted result is a 267× orchestration speedup with zero false positives across the benchmark corpus. This is the throughput-axis evidence at the substrate's scale.

6.4 Model-Derived Scaling Projections (Not Measured Data)

No multi-scale measurement grid was operated for this work. To make the formal model of Sections 4 and 5 concrete and falsifiable, this section presents the model's quantitative predictions at three illustrative configurations: 8 nodes / 3 agents, 32 nodes / 12 agents, and 128 nodes / 48 agents. Every figure below is a projection derived from the closed-form expressions of Section 4.2 — not a measurement. The exponent $\alpha \approx 0.79$ used in the throughput and cost projections is an illustrative value within the (0,1) bound established in Section 4.2.2; it is not fitted to measured multi-scale data, and measuring the actual α is a

primary purpose of ICS-Bench.

Table 2. Model-Projected ICS Metrics at Three Illustrative Scales (projections of the Section 4.2 closed-form model — not measurements)

Metric	Small (8N / 3A)	Medium (32N / 12A)	Large (128N / 48A)	Projected S→L change
MTBF (hours)	4.2	31	187	+44.5×
Throughput (ops/sec)	840	4,200	23,800	+28.3×
Latency variance (ms ²)	142	38	9	−15.8×
Cost per 1K ops (USD)	\$0.84	\$0.19	\$0.047	−17.9×

If these projections are correct, they imply an 18× cost reduction and a 28× throughput improvement between the modeled scales. The figures quoted in Sections 7.2 and 9 are these projections and must not be cited as measurements. Any citation of this paper as demonstrating measured multi-scale ICS would be a misreading of this section.

6.5 Discussion of Confounds

ICS is not claimed to be universal across all workload types. The model predicts the phenomenon will be strongest for workloads that are routing-heavy (benefiting from route diversity), agent-parallel (benefiting from Executor scaling), and stateless per operation (enabling full horizontal scalability); the substrate's AI inference routing is the workload class for which the strongest ICS signature is predicted.

ICS is expected to weaken measurably for strongly sequential workloads, where Amdahl's Law constraints on the serial fraction limit the benefit of additional Executor agents. Similarly, externally rate-limited workloads (where throughput is constrained by an upstream API rate limit rather than internal system capacity) are not expected to exhibit throughput ICS, though the stability and cost projections may persist. Finally, stateful workloads requiring per-session state migration impose context-graph write contention that partially offsets the shared-context memory efficiency benefit. Each of these expectations is a testable hypothesis for ICS-Bench, not a receipted result.

7. Implications and Applications 7.1 Infrastructure Design Philosophy ICS inverts the classical "start small and keep it simple" heuristic for DSOS/CAI architectures. Traditional wisdom holds that complexity should be added only when required by demonstrated need, because each increment of complexity incurs overhead. Under ICS, this logic produces the worst possible outcome: an underscaled DSOS/CAI system is simultaneously more expensive per operation, less stable, less deterministic, and lower in throughput than a properly scaled one. The revised design principle is: design for complexity headroom from day one. For DSOS/CAI systems, the deployment should be sized to at least kinflection from initial launch. The marginal cost of additional nodes and agents at deployment time is far lower than the operational cost differential of running in the fragile zone below kmin. This is a non-trivial organizational shift: it requires capital expenditure commitments at launch that traditional "start lean" methodologies would defer. The ICS cost model (Section 7.2) provides the quantitative justification for this investment. 7.2 Cost Modeling for DSOS/CAI Systems Traditional infrastructure cost models project total cost of ownership (TCO) as approximately linear with resource count:

$TCO_{\text{classical}}(\kappa) \approx \text{Cunit} \cdot \kappa$ This model is incorrect for DSOS/CAI systems exhibiting ICS. The correct ICS-aware TCO model accounts for the super-linear growth of operational output with complexity: $TCO_{\text{ICS}}(\kappa) = \text{Cinfra}(\kappa) / \text{Ops}(\kappa)$ where $\text{Ops}(\kappa) = \Omega(\kappa^\alpha)$ Since $\text{Cinfra}(\kappa)$ grows approximately linearly with κ (linear resource cost), while $\text{Ops}(\kappa)$ grows as κ^α with $\alpha \approx 0.79$ in the illustrative model calibration of Section 6.4 plus overhead reduction, the peroperation TCO falls with scale. Engineering teams that use classical linear TCO projections will systematically overestimate the cost of larger DSOS/CAI deployments and underestimate the cost of smaller ones.

7.3 Comparison to Classical Distributed Systems

It is instructive to contrast ICS-exhibiting DSOS/CAI architectures against the two dominant classical paradigms. Microservices architectures decompose functionality into many small, independently deployable services. While this achieves modularity and independent scalability, the proliferation of services increases the failure surface: each interservice call is a potential failure point, and service mesh complexity grows with service count. The topology is typically fixed and engineered, not self-organizing; there is no equivalent to DSOS's self-healing route tables. Microservices do not exhibit ICS. Monolithic architectures centralize complexity in a single codebase and process. Performance degrades with feature additions because the serial fraction grows, coordinating resources become bottlenecks, and deployment risk increases with every change. Monoliths are the canonical case of classical complexity-penalty behavior. The differentiating factor in DSOS/CAI is the combination of self-organizing topology, role-differentiated agent specialization, and shared append-only context. These three properties together convert complexity from a penalty into an asset—a combination that neither microservices nor monoliths achieve.

7.4 Roadmap for Future Validation

The results presented in this paper, while internally consistent and theoretically grounded, constitute a single experimental environment. The following research roadmap is proposed for community validation:

- **ICS-Bench:** A formal, open-source benchmark suite providing standardized workloads (routing-heavy, agent-parallel, mixed-stateful), measurement protocols, and reporting templates. ICS-Bench should enable independent replication of ICS results across diverse cloud environments and organizational deployments.
- **Heterogeneous node/agent populations:** This paper assumes homogeneous node and agent types within each category. Real-world deployments often involve heterogeneous hardware (varying CPU, memory, network bandwidth). The ICS model should be extended to account for heterogeneous failure rates and capacities.
- **Empirical characterization of k_{max} :** The ICS breakdown threshold is theoretically predicted but not empirically characterized in this work. Controlled experiments pushing beyond current deployment scales are needed to map the breakdown zone and determine whether it is workload-dependent.
- **Correlated failure modeling:** The current model assumes node failure independence. Correlated failure events (datacenter power outages, network partition events, software bugs affecting multiple nodes simultaneously) violate this assumption. Extensions incorporating correlated failure models are an important priority.

8. Limitations and Threats to Validity

We acknowledge the following limitations of this work:

7. **Bounded operational range.** ICS is formally defined and predicted only within $[k_{\text{min}}, k_{\text{max}}]$. Behavior beyond k_{max} is not characterized; extrapolation beyond the modeled range is not supported by any data in this work.
8. **Single-scale receipted evidence and projected multi-scale figures.** All receipted operational data was collected from a single production substrate (25 nodes, 10 agents) operated by the author. No multi-scale deployment grid (such as the 8/32/128-node configurations of Section 6.4) was operated for this work; every multi-scale figure in this paper is a model-derived projection, not a measurement. External

validity — whether ICS results replicate on alternative cloud providers, hardware configurations, and organizational workloads — requires broader independent replication. The author strongly encourages this replication. 9. Infrastructure-only cost accounting. The cost figures reported in Section 6.5 cover only cloud infrastructure costs (compute, memory, network egress). They do not include engineering labor costs for deployment, maintenance, or incident response. Engineering overhead may exhibit different scaling properties and could offset some cost-efficiency gains at certain scales, particularly in early deployment phases. 10. Node failure independence assumption. The mathematical derivations in Section 4.2 assume that individual node failures are mutually independent. This assumption is violated in the presence of correlated failure events (shared hardware, software vulnerabilities affecting multiple nodes, network partition events). In such scenarios, the super-exponential reliability improvement may not materialize, and the practical MTBF gains would be reduced. The Raft consensus mechanism provides some protection against correlated failures but does not eliminate them. 11. Workload generalizability. The three workload types tested were specifically selected to be representative of DSOS/CAI use cases. They are likely not representative of all distributed systems workloads, and ICS should not be expected to hold for arbitrary workloads in arbitrary distributed architectures.

9. Conclusion This paper has introduced, formally defined, mathematically derived, architecturally explained, and supported with receipted operational evidence a new phenomenon in distributed systems: Inverse Complexity Scaling (ICS). In DSOS and CAI architectures, the classical assumption that complexity degrades performance is violated in a consistent and quantifiable way. As node count, agent count, and route density increase within the operational range $[k_{min}, k_{max}]$, all four key performance metrics—stability, throughput, determinism, and cost-efficiency—improve simultaneously. The formal definition of ICS (Definition 1) provides a precise, testable characterization of this phenomenon. The mathematical derivations show that ICS arises inevitably from the combination of super-exponential fault absorption through route redundancy, order-statistics minimization of expected latency via route diversity, linear load distribution across agent populations, and fixed-cost amortization across growing operation counts. These are not incidental properties of a particular implementation; they are structural consequences of self-organizing topology, role-differentiated agents, and shared context graphs. The receipted operational evidence is consistent with the ICS model: zero false-positive heals and zero sub-threshold executions across a 200-row production incident annex under frozen per-playbook confidence thresholds, an 84.5% auto-heal ratio on a 400-injection chaos run against an 88.5% historical baseline, and a 267× orchestration benchmark speedup on the 25-node, 10-agent production substrate. The multi-scale figures (up to 18× cost reduction, 28× throughput improvement, 44.5× MTBF improvement, and 15.8× variance reduction between modeled scales) are model projections, not measurements; converting them into measured evidence is precisely the purpose of ICS-Bench. Together, the receipted evidence and the model support ICS and challenge the prevailing engineering wisdom that DSOS/CAI-class systems should be sized conservatively. The practical implication is direct: for self-organizing, agent-parallel distributed systems, under-scaling is the more expensive mistake. Infrastructure teams and system architects should adopt ICS-aware cost models and design for complexity headroom from initial deployment. The classical "start small" heuristic, appropriate for most system architectures, is contraindicated for systems exhibiting ICS. We invite the systems and AI research

community to independently validate these findings through the proposed ICS-Bench benchmark suite, to extend the formal model to heterogeneous and correlated-failure settings, and to collaboratively establish the empirical boundaries of ICS across diverse deployment environments.

The phenomenon reported here is sufficiently significant that it warrants broad, rigorous community engagement. In an era where distributed AI orchestration is becoming a foundational infrastructure concern, ICS offers a principled, mathematically grounded basis for infrastructure design decisions. We believe this foundation, once widely validated, will meaningfully reshape how self-organizing AI systems are built, scaled, and costed.

References ● [1] F. P. Brooks Jr., "No Silver Bullet: Essence and Accidents of Software Engineering," *Computer*, vol. 20, no. 4, pp. 10-19, Apr. 1987. doi: 10.1109/MC.1987.1663532. ● [2] A. S. Tanenbaum and M. Van Steen, *Distributed Systems: Principles and Paradigms*, 3rd ed. Upper Saddle River, NJ, USA: Prentice Hall, 2017. ● [3] G. M. Amdahl, "Validity of the Single Processor Approach to Achieving Large Scale Computing Capabilities," in *Proc. AFIPS Spring Joint Computer Conf.*, Atlantic City, NJ, USA, Apr. 1967, pp. 483-485. doi: 10.1145/1465482.1465560. ● [4] J. L. Gustafson, "Reevaluating Amdahl's Law," *Commun. ACM*, vol. 31, no. 5, pp. 532-533, May 1988. doi: 10.1145/42411.42415. ● [5] M. E. Conway, "How Do Committees Invent?" *Datamation*, vol. 14, no. 4, pp. 28-31, Apr. 1968. ● [6] A. Avizienis, J.-C. Laprie, B. Randell, and C. Landwehr, "Basic Concepts and Taxonomy of Dependable and Secure Computing," *IEEE Trans. Dependable Secure Comput.*, vol. 1, no. 1, pp. 11-33, Jan. 2004. doi: 10.1109/TDSC.2004.2. ● [7] L. Lamport, R. Shostak, and M. Pease, "The Byzantine Generals Problem," *ACM Trans. Program. Lang. Syst.*, vol. 4, no. 3, pp. 382-401, Jul. 1982. doi: 10.1145/357172.357176. ● [8] D. Ongaro and J. Ousterhout, "In Search of an Understandable Consensus Algorithm," in *Proc. USENIX Annual Technical Conf. (USENIX ATC)*, Philadelphia, PA, USA, Jun. 2014, pp. 305-319.

- [9] R. Axelrod and M. D. Cohen, *Harnessing Complexity: Organizational Implications of a Scientific Frontier*. New York, NY, USA: Free Press, 1999. ● [10] M. Dorigo and T. Stützle, *Ant Colony Optimization*. Cambridge, MA, USA: MIT Press, 2004. ● [11] A. Demers, D. Greene, C. Hauser, W. Irish, J. Larson, S. Shenker, H. Sturgis, D. Swinehart, and D. Terry, "Epidemic Algorithms for Replicated Database Maintenance," in *Proc. 6th ACM Symposium on Principles of Distributed Computing (PODC)*, Vancouver, BC, Canada, Aug. 1987, pp. 1-12. doi: 10.1145/41840.41841. ● [12] M. Wooldridge, *An Introduction to MultiAgent Systems*, 2nd ed. Chichester, UK: John Wiley & Sons, 2009. ● [13] J. Dean and S. Ghemawat, "MapReduce: Simplified Data Processing on Large Clusters," *Commun. ACM*, vol. 51, no. 1, pp. 107-113, Jan. 2008. doi: 10.1145/1327452.1327492. ● [14] L. Lamport, "Paxos Made Simple," *ACM SIGACT News*, vol. 32, no. 4, pp. 18-25, Dec. 2001. ● [15] M. Castro and B. Liskov, "Practical Byzantine Fault Tolerance," in *Proc. 3rd USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, New Orleans, LA, USA, Feb. 1999, pp. 173-186. ● [16] S. Shenker, "Fundamental Design Issues for the Future Internet," *IEEE J. Select. Areas Commun.*, vol. 13, no. 7, pp. 1176-1188, Sep. 1995. doi: 10.1109/49.414653. ● [17] H. Garcia-Molina, J. D. Ullman, and J. Widom, *Database Systems: The Complete Book*, 2nd ed. Upper Saddle River, NJ, USA: Pearson Prentice Hall, 2009. ● [18] Leon Calvin Long II, "DSOS Architecture Reference: Self-Healing Route Tables and Local Consensus in Distributed Orchestration Systems," *DSOS-CAI Research Technical Report TR-2025-01*, Spicewood, TX, USA, 2025. [Internal preprint; community release pending ICS-Bench publication]. Inverse Complexity Scaling in DSOS and CAI Systems — Leon

Calvin Long II (DSOS-CAI Research) — September 2026 — Version 2 (Amended) —
Spicewood, TX, USA

This document is released as a research preprint (Version 2, amended). Receipted results are from the author's production substrate; all multi-scale figures are model-derived projections, labeled as such. Independent replication encouraged. | ICS-Bench proposal: open to community contribution.